

SEA 2018 report

Contents

1	Experimental setup	1
1.1	Time measurements	1
2	WL tests	1
2.1	Input data	1
2.2	Double rank versus single rank	2
2.3	Lazy versus nonlazy	3
2.4	Double rank and fail versus double rank	4
2.5	Double rank, Lazy, and Rank and fail versus Single rank, nonlazy, and nonfail	5
2.6	Maxrep	6
3	Optimizations on parent operations	7
3.1	Input data	7
3.2	LCA versus parent sequence	8
4	Parallelization	9
5	Genome tests	12
5.1	Figure 4	12
5.2	Figure 5	13
5.3	parallel on real data	13
6	Range queries	13

1 Experimental setup

1.1 Time measurements

We run a program k times with and without the optimization and recorded the sum construction time of the MS and RUNS vectors. The plots report (median, with quartile ranges) of the speedup of each optimized time t_i^{opt} relative to the average non-optimized time in the construction time of the MS vector. In other words

$$d^{(i)} = \frac{\bar{t}_{\text{non_opt}}}{t_{\text{opt}}^{(i)}}$$

with $\bar{t}_{\text{non_opt}} = 1/n \sum t_{\text{non_opt}}^{(i)}$, and $i = 1, \dots, k$.

The boxplots report the raw times.

2 WL tests

2.1 Input data


We perform tests on the Weiner Link optimizations on 4 types of input.

- Index string with repeats, query string random (code: `rep_dis`)


- Index string with repeats, query string similar to index (code: `rep_sim`)
- Index string random, query string random (code: `rnd_dis`)
- Index string random, query string similar to index (code: `rnd_sim`)

Further, we generate all of the above input data for two alphabet sizes: $|\Sigma_1| = 4$ and $|\Sigma_2| = 20$. For all input types, the index string is of length 100MB and the query 500KB.


2.2 Double rank versus single rank




2.3 Lazy versus nonlazy



2.4 Double rank and fail versus double rank




2.5 Double rank, Lazy, and Rank and fail versus Single rank, nonlazy, and nonfail




2.6 Maxrep

2.6.1 Vanilla Maxrep vs. non-maxrep



2.6.2 Rank&check Maxrep versus Vanilla Maxrep



3 Optimizations on parent operations

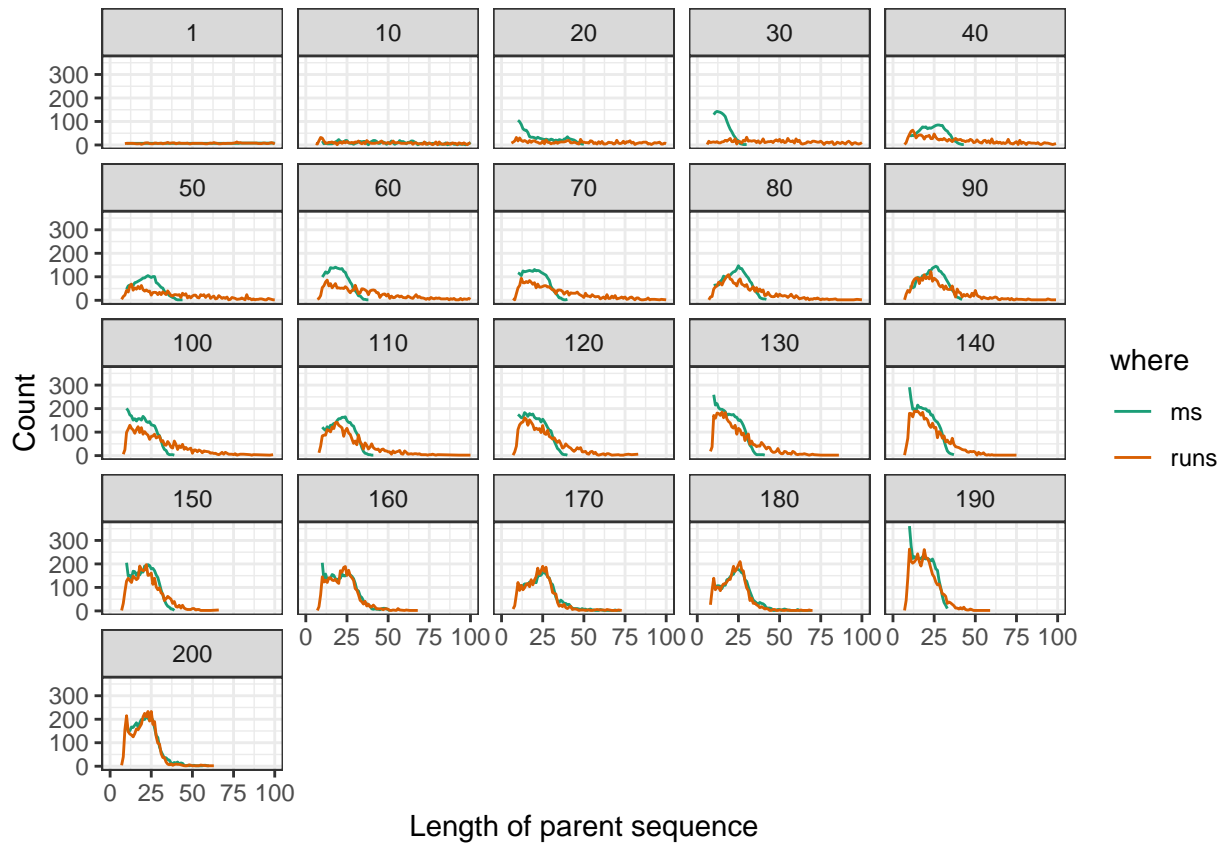
3.1 Input data

We generate the index input string with repetitions as follows. We generate a random seed block b of length 200. Next, we generate blocks of the same length b_k by introducing k mutations on b . The index string of length 10MB is $b \circ b_k^{(1)} \dots \circ b_k^{(4999)}$.

The query string is obtained as a concatenation of labels from nodes of the suffix tree of s . We select nodes with node depth of at least 10 and string length at most 170 for a total string length of 103KB. We separate the labels with a sentinel character that does not appear in s .

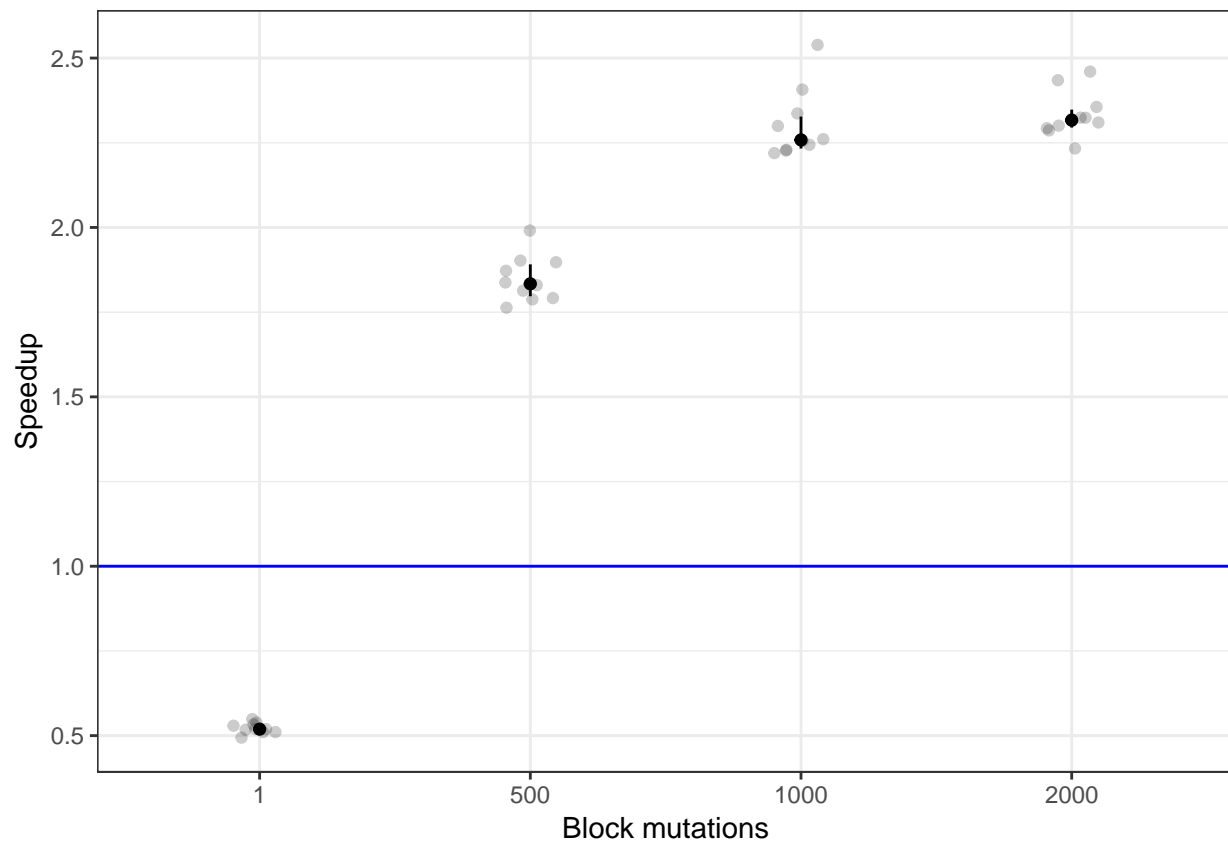
Furthermore, we perform experiments for various choices of $1 \geq k \geq |b|$.

The plot below shows a histogram of the length of consecutive parent operations. This quantity is important since the speedup of this optimization is proportional to the length of sequence of parent operations. Importantly, the optimization might not even be beneficial if the length of the sequence of parent operations is less than 3.





3.2 LCA versus parent sequence



```
## # A tibble: 40 x 5
## # Groups:   ntrial [10]
##   ntrial k      lca pseq value
##   <int> <fct> <dbl> <dbl> <dbl>
## 1     1 1      887  459 0.517
## 2     1 1000   207  476 2.30
## 3     1 2000   213  487 2.29
## 4     1 500    225  408 1.81
## 5     2 1      863  448 0.519
## 6     2 1000   209  469 2.24
## 7     2 2000   216  502 2.32
## 8     2 500    234  430 1.84
## 9     3 1      861  459 0.533
## 10    3 1000   205  479 2.34
## # ... with 30 more rows
```

4 Parallelization


```
## # A tibble: 240 x 5
##   inp_type nthreads section action    avg
##   <chr>      <int> <chr>   <chr>  <dbl>
## 1 rnd_dis      1 comp    total  3373
## 2 rnd_dis      1 ms      build  1528
## 3 rnd_dis      1 ms      total  1633
## 4 rnd_dis      1 runs    build  1593
## 5 rnd_dis      1 runs    correct 14
## 6 rnd_dis      1 runs    total  1739
## 7 rnd_dis      2 comp    total  1689
## 8 rnd_dis      2 ms      build   745
## 9 rnd_dis      2 ms      total   824
## 10 rnd_dis     2 runs    build   769
## # ... with 230 more rows
```






5 Genome tests

5.1 Figure 4



5.2 Figure 5



5.3 parallel on real data

6 Range queries

```
## # A tibble: 15 x 5
##   block_size range_size nqueries time_ms time_per_query
##   <dbl>      <dbl>      <dbl>   <dbl>      <dbl>
## 1         0 20000000         1     56         56
## 2         0 40000000         1    177        177
## 3         0 60000000         1    296        296
## 4         0 100000000        1    537        537
## 5         0 200000000        1    947        947
## 6         4 20000000 1000000    328    0.000328
## 7         4 40000000 1000000    352    0.000352
## 8         4 60000000 1000000    372    0.000372
## 9         4 100000000 1000000    391    0.000391
## 10        4 200000000 1000000    391    0.000391
## 11      1024 20000000 1000000    327    0.000327
## 12      1024 40000000 1000000    344    0.000344
## 13      1024 60000000 1000000    366    0.000366
## 14      1024 100000000 1000000    387    0.000387
## 15      1024 200000000 1000000    392    0.000392
```

